

George Mason University
Learning From Data Fall 2024
Computer Exercise #8 Part 2

Assigned: November 20, 2024

Due Date: December 06, 2024

Computer Exercise 8.2 (Image Classification):

Now that you have some familiarity with CNNs, you are now tasked to design an image classifier using the dataset provided with this exercise. Your goal is to learn a classifier that you believe will perform the best on a test set that will be used to evaluate your classifier.

- (a) The dataset contains labeled images of dogs that are stored in an `.npz` file, a compressed archive format used by NumPy. There are 3,132 color images scraped from the internet and resized to 224×224 to align with the size of the images in the ImageNet dataset. Each image belongs to one of twenty different classes, with class labels zero through 19. The name of the file is `train_data_224.npz` and is 471.5MB in size.
- (b) The dataset may be read and put into two arrays, `images` and `labels`, as follows:

```
# Load the .npz file
data = np.load('train_data_224.npz')
# Access the arrays stored in the .npz file
images = data['images'] # Access the images array
labels = data['labels'] # Access the labels array
```

- (c) Before proceeding, you may wish to become familiar with the dataset.

- (a) Look at some images,

```
plt.imshow(images[0], interpolation='nearest')
plt.show()
```

- (b) Examine the number of samples in each class

```
unique_labels, counts = np.unique(labels, return_counts=True)
# Print the counts
for label, count in zip(unique_labels, counts):
    print(f"Label {label}: {count}")
```

and determine whether or not the dataset is balanced, and if there are enough images in each class necessary for the type of classifier you will be learning.

- (d) Consider what data preprocessing tasks you think are appropriate for this problem.
- (e) You may wish to consider using `tf.keras.preprocessing.image.ImageDataGenerator`.
- (f) Your *classic*-style classifier of convolution and pooling layers followed by one or more fully connected layers may not perform as well as you would like, so you may wish to consider using transfer learning.

What you are to submit

Your report should be included in the PDF file you will be submitting for Part 1. For Part 2, you should include the following:

1. A brief discussion of the types of classifiers you considered and how well they worked.
2. A detailed description of your final classifier, the architecture, the preprocessing and how the training is done.
3. A learning curve, the classification accuracy, and a confusion matrix.

Don't forget to include a code cell in your program with the following:

```
model.save("model.NNNN") # Saves both architecture and weights
```

This will save the architecture and parameters of your classifier and allow it to be evaluated on an evaluation dataset. Replace NNNN with the **last four digits of your G-Number**. You **must** submit this file along with your writeup. Failure to do so will result in zero points for part 2. You **must also** make sure that this file can be read by an evaluation program and properly evaluate the classifier on an evaluation dataset. To ensure that your saved file will work, the Jupyter notebook used to evaluate your classifier is provided along with a sample evaluation dataset. Again, if the evaluation program is unable to read in your model architecture and parameters to evaluate a dataset, then you will receive a grade of zero for part 2.